

Joseph Seaton

Shader Compositor

Computer Science Tripos, Part II

Fitzwilliam College

May 15, 2012

Proforma

| | |
|---------------------|----------------------------|
| Name: | Joseph Seaton |
| College: | Fitzwilliam College |
| Project Title: | Shader Compositor |
| Word Count: | About 9500 |
| Project Originator: | Christian Richardt |
| Supervisor: | Christian Richardt |

Original aims of the project

The aim of this project was to produce a novel user interface for easy composition and testing of a certain class of graphical effects known as shaders. The user interface was to be highly interactive and allow easy modification of the code for the effects themselves, the manner in which they are composed and, via graphical means, their associated parameters. Previews of these changes were to be provided as quickly as possible; therefore the project also included a number of optimisations.

Work completed

The core of the project has been completed and works satisfactorily. Shaders can be created and a tree of shaders can be specified using JavaScript, the result of which is shown in-browser. An interface for modifying shader parameters is correctly generated for all types of parameter that have been tested, and respects some annotations. A number of optimisations have been applied, including elimination of unnecessary shader recompilation and pipeline

re-generation. The project is also now capable of reusing framebuffer objects between pipeline changes where possible.

Special difficulties

None.

Declaration

I, Joseph Seaton of Fitzwilliam College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Signed

Date

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 9 |
| 1.1 | Motivation | 9 |
| 1.2 | Brief introduction to shaders | 10 |
| 1.2.1 | A look at some OpenGL calls | 10 |
| 2 | Preparation | 13 |
| 2.1 | Requirements analysis | 13 |
| 2.2 | Preparatory learning | 14 |
| 2.2.1 | A short introduction to OpenGL | 14 |
| 2.3 | Programming language and libraries | 14 |
| 2.4 | Development environment | 18 |
| 2.5 | Software development process | 20 |
| 3 | Implementation | 21 |
| 3.1 | GLSL Parser | 23 |
| 3.1.1 | Grammar conversion | 24 |
| 3.1.2 | Preprocessor | 24 |
| 3.1.3 | Annotations | 25 |
| 3.1.4 | Struct parsing | 26 |
| 3.2 | Pipeline specification | 28 |
| 3.3 | Parameter UI generation | 31 |
| 3.4 | Pipeline generation | 33 |
| 3.5 | Rendering | 34 |
| 3.6 | Optimisations | 35 |
| 3.6.1 | GLOW early initialization | 35 |
| 3.6.2 | Modified shader detection | 35 |
| 3.6.3 | Parameter value propagation | 36 |
| 3.6.4 | GLOW cache | 36 |
| 3.6.5 | FBO reuse | 36 |

| | | |
|----------|--------------------------------|-----------|
| 3.7 | User Interface | 37 |
| 3.7.1 | Editor | 37 |
| 3.7.2 | Layout | 37 |
| 3.7.3 | Improved Layout | 37 |
| 4 | Evaluation | 41 |
| 4.1 | Parser correctness | 41 |
| 4.2 | Pipeline correctness | 44 |
| 4.3 | Performance | 46 |
| 5 | Conclusion | 49 |
| 5.1 | Main Results | 49 |
| 5.2 | Lessons learned | 49 |
| 5.3 | Further work | 51 |
| | Bibliography | 53 |
| A | Project Proposal | 55 |

Chapter 1

Introduction

My project provides a novel user interface for easily composing multiple graphical effects known as shaders together, in which individual shader parameters can be easily modified and the results viewed immediately. Furthermore the interface is sufficiently optimised to provide a low round-trip time between making a change and seeing the effect of this change. I have successfully implemented the core of my project, as well as a number of proposed extensions and a few optimisations not initially proposed.

1.1 Motivation

Working with raw OpenGL is undeniably messy. OpenGL is a verbose and complicated API. When developing shaders, developers tend not to be interested in writing large amounts of OpenGL API code just to tweak shader parameters or forward the texture output of one shader to the input of another. But this is exactly what developers have to do. Given that many graphics applications are written in compiled languages like C and C++, this significantly increases the length of a development cycle. As a large part of shader development constitutes small tweaks to (often subjectively) improve output, this is not ideal.

For students learning about shaders, the OpenGL API is often an intimidating and confusing mess that gets in the way of understanding a fundamentally quite straightforward concept. Listing 1 shows some example OpenGL calls.

This project aims to provide a way of developing shaders that is both useful for developers, and may be used as a tool for students learning about shaders for the first time.

1.2 Brief introduction to shaders

In their modern incarnation, OpenGL shaders are quite a simple concept to understand at a high level. A shader is a simple program — ‘simple’ referring to certain constraints we shall ignore for now — written in a dataflow style. Such programs take in one item of data at a time, for example a vertex, and output another item of data — the new vertex location, or a pixel value.

In a modern OpenGL implementation, there are two commonly used types of shader (and a few more which will not be discussed). These are vertex shaders and fragment shaders. Vertex shaders operate on individual vertices, transforming the position and also attaching information such as a colour. Fragment shaders roughly speaking generate pixel values. Information can be passed from vertex shaders to fragment shaders using varying parameters, the values of which are interpolated from those of the surrounding vertices.

1.2.1 A look at some OpenGL calls

As discussed above, the use of shaders sounds simple. However, as we can see from Listing 1, a short excerpt of the OpenGL calls from an actual WebGL program (author?) [16], the paraphernalia of mostly boilerplate OpenGL calls required for compiling shader programs, specifying parameters and so on can result in a very large codebase for even the simplest program.

Listing 1 Example WebGL calls

```
glBindFramebuffer(GL_FRAMEBUFFER, 0);
glUseProgram(ProgramName[IMAGE_2D]);
glUniformMatrix4fv(UniformMVP[IMAGE_2D], 1, GL_FALSE, &MVP[0][0]);
glUniform1i(UniformDiffuse, 0);

glActiveTexture(GL_TEXTURE0);
glBindTexture(GL_TEXTURE_2D_ARRAY, TextureColorbufferName);
glBindSampler(0, SamplerName);

glBindVertexArray(VertexArrayName[IMAGE_2D]);
glBindBuffer(GL_ELEMENT_ARRAY_BUFFER, BufferName[BUFFER_ELEMENT]);
```

Chapter 2

Preparation

Before actually embarking on the project, there were a few decisions to be made, as discussed in this chapter. Firstly, project planning...[TODO]

2.1 Requirements analysis

Before beginning the project, it was important to establish the general architecture. In 2.1 we can see a summary of the main objectives of the project. [Is this reasonable detail?]

In order to properly plan the development of the project, it was necessary to construct a graph of the dependencies between parts of the project. Priority could then be given to those parts of the project that were necessary for other parts of the core. Figure 2.1 shows the general dependencies between larger parts of the project.




Figure 2.1: Project dependencies.

| Goal | Priority |
|--------------------------------------|----------|
| Simple interface | High |
| Parameter detection and presentation | High |
| Pipeline generation | High |
| Pipeline rendering | High |
| Syntax highlighting | Medium |
| Optimisations | Medium |
| Annotations | Medium |
| Additional UI work | Low |
| Animation | Low |

Table 2.1: Main goals

2.2 Preparatory learning

2.2.1 A short introduction to OpenGL

OpenGL is a hardware-, operating-system and vendor-independent graphics API. It has a stateful, client-server architecture, in which the client, the OpenGL library sends commands — OpenGL API calls — to the server, the actual OpenGL driver. These commands alter the current state, specifying for example that the following commands will specify a polygon. These commands enter a pipeline, a very simplified version of which is shown in Figure [?]. The vertex and fragment shaders discussed in Section ?? are contained within the per-vertex and per-fragment operation blocks respectively. Of particular concern to for this project is the final stage in this pipeline, in which fragments are rasterized and (following some other processing for example z-buffering) written to a framebuffer. A framebuffer object or FBO contains several buffers including a colour buffer, which may accessed as a texture by other stages.

2.3 Programming language and libraries

WebGL

WebGL is a very recently developed API allowing web-based applications access to OpenGL ES contexts, via JavaScript (**author?**) [8]. I decided that the use of web-based technologies would help lower the barrier for use of this

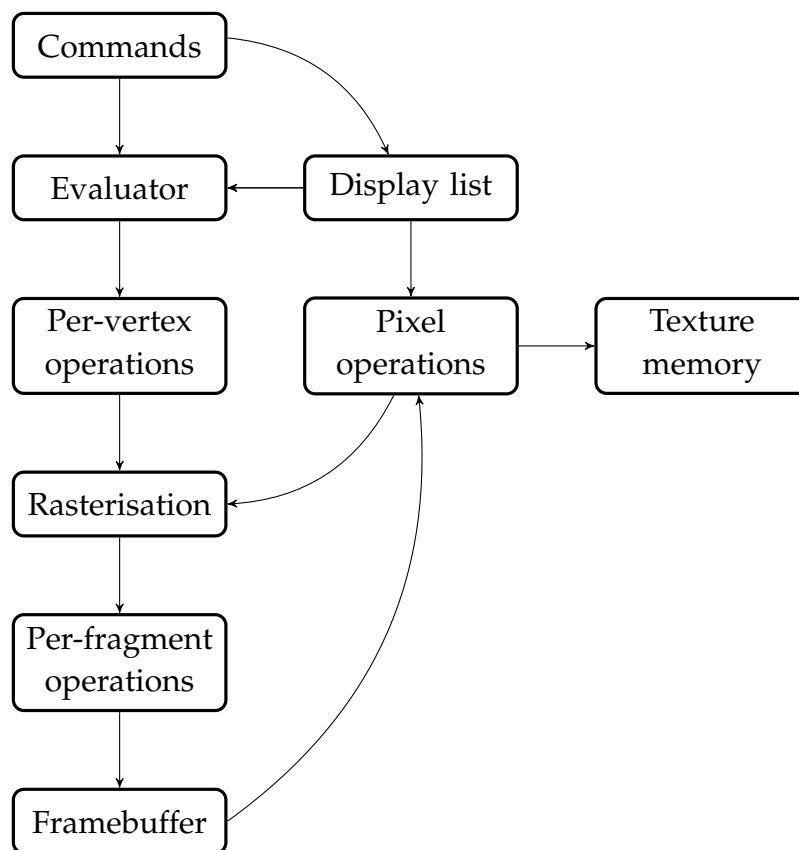


Figure 2.2: Simplified OpenGL pipeline.

project, and therefore shader development in general, since a user needs only visit the correct page using a modern web browser. Using WebGL has the added advantage of requiring less familiarisation given my prior experience with JavaScript. WebGL uses essentially the same API as OpenGL ES for C **(author?)** [8].

OpenGL and GLSL

GLSL is the only choice of shading language for WebGL, and is the standard high-level shading language of OpenGL. Syntactically it is very similar to C, and is intended to be written by humans. Alternatively for a project not using WebGL, I could have chosen to use Cg, a language with similar goals but the ability to output OpenGL and DirectX compatible shader programs. However WebGL does not support Cg and does not appear to be as well known.

Prior to beginning this project, I had very little knowledge of OpenGL barring a small amount of personal experience writing toy programs in my spare time. As such, I needed to gain a better understanding of OpenGL, and some knowledge of GLSL before I could begin this project. For this purpose, I obtained a copy of the famous ‘Red Book’ **(author?)** [12]. I also read through the GLSL ES specification provided by the Khronos Group **(author?)** [14], which discusses the syntax of GLSL in detail.

JavaScript

JavaScript is the natural choice for any client-side web development work, and it is currently unavoidable that any WebGL project will have to at least output JavaScript at some point. While JavaScript has a less than stellar image due to implementation inconsistencies between browsers, the language itself is quite sophisticated, featuring first-class functions, closures, and a prototype-based object model.

While, like many people, I already have some experience of JavaScript programming from web development work in the past, I thought it would be wise to gain a broader understanding of JavaScript before embarking on a more complex project such as this one. I found particularly useful *Eloquent Javascript* **(author?)** [10] and *Higher Order Programming [in JavaScript]* **(author?)** [15], both of which are unusual among texts on JavaScript in being willing to discuss

the more sophisticated functional aspects of JavaScript, and their relation to JavaScript's own unusual prototype-based object system.

jQuery

jQuery is a general utility library for JavaScript, which provides many helper functions not included in web browsers, and also greatly simplified DOM (Document Object Model) manipulation. In Listing 2 we see the difference in concision with and without jQuery for some simple common tasks.

Listing 2 Comparison of plain JavaScript with jQuery for adding a 'click' event handler to all links ('a')

```
// JavaScript
[].forEach.call(document.querySelectorAll("a"), function(el) {
    el.addEventListener("click", function() {
        ...
    });
});

// jQuery
$("a").click(function() {
    ...
});
```

Testing framework

There are many unit testing frameworks available for JavaScript, of varying levels of complexity. Many of these frameworks are designed with very large projects and extensive test suites in mind. For my project, I considered these to be overly complicated to work with. Eventually, I settled on QUnit^{??}, a very simple unit testing framework that is part of jQuery, and it's port to node.js (a fork of the original code) ^{??}.

WebGL toolkit

Given that OpenGL and therefore WebGL itself can be quite awkward to work with, I decided to use an existing library to abstract away some of the boil-

erplate code irrelevant to my project. However, there are currently many libraries available claiming to do exactly this, of varying degrees of completeness and varying levels of abstraction. Initially, I settled on three.js (**author?**) [5] on the basis of its popularity and high rate of development. However, it soon became apparent that the level of abstraction provided by three.js actually made the task more difficult, given that interacting with FBOs was non-obvious. After some more searching, I came across GLOW (**author?**) [2], a toolkit specifically designed to make working with shaders simple, but otherwise providing little abstraction.

Editor

One important component of the project is to provide good GLSL syntax highlighting. There currently exist many web-based editors offering some degree of syntax highlighting, but I chose to use CodeMirror (**author?**) [1], given its active development and the presence of a modular way to add support for highlighting new grammars.

jQuery UI

jQuery UI is a user interface library built on top of jQuery. It provides a number of common, basic widgets including a simple tabbing interface. Using a pre-existing library is especially useful for user interface work in JavaScript, since widgets will have been tested against multiple browsers, a time-consuming and tricky process.

2.4 Development environment

Emacs

Surprisingly for such an old editor, with the correct extensions, Emacs is very well equipped for modern JavaScript development. While there are other editors available that could claim most or all of the features these extensions add, my previous experience with Emacs makes this a favourable choice. The extensions used add auto-completion, syntax checking, and an inline node.js (see below) console (**author?**) [11].

node.js

node.js is a JavaScript platform built on top of Google's V8 JavaScript engine (author?) [6]. It provides, among other things, a command line interface, including a read, evaluate, print loop (or REPL). This provides a fast, lightweight way to test non-browser-bound code using traditional command line tools, can be easily interfaced with git and cron, and can easily write to files. All of this can be done without requiring messy extra code to interact with a web browser or server-side code to report back statistics.

Google Chrome

WebGL naturally requires the use of a web browser. I have found that Google Chrome provided the most reliable WebGL implementation on my development machine, which runs Linux on an ATI HD 6850 graphics card. Since WebGL is a relatively recent development that has only recently gained browser support, WebGL can still be unstable and buggy in some circumstances — even using Google Chrome. Chrome also includes a set of 'developer tools' including a JavaScript console

WebGL inspector

WebGL inspector is a Google Chrome extension intended to provide inline information on current WebGL contexts, giving easy access to state information including currently allocated textures. While this tool would be very useful for debugging, I found the tool to be insufficiently stable on my operating system (Linux) to use reliably. This meant that most inspection of OpenGL state had to be performed manually.

Version control and backup strategy

All project files, including the dissertation, were placed under git version control. Git is a distributed version control system, thus it was possible to maintain multiple complete copies of the repository in different locations. Git was then configured to push to multiple remote repositories in different locations: an external harddrive, my PWF account, my SRCF account, my web hosting

account, and GitHub ??). While this could be considered excessive, the use of SSH keys meant that post-setup, this required no extra effort on my part. Since I am conscious that git is capable of history rewriting, I also configured a regular cron job to make regular compressed backups of my current repository using `git bundle`.

2.5 Software development process

WebGL is still quite a new technology, and as such the libraries surrounding it are not yet mature. Given that I was also not overly familiar with OpenGL at the beginning of the project, I deemed it sensible to apply an iterative software development model, to allow the requirements to be updated after prototyping.

I decided to use an evolutionary software development model. This model employs rapid prototyping, enabling fast experimentation and evaluation of the feasibility of a given route. The model also incorporates feedback from development into the specification.

Test-driven development was also employed for some parts of the project, most prominently the parser. The use of unit testing enables early, automated detection of regressions, making it easy to spot bugs that might otherwise go unnoticed. Git pre-commit hooks were used to enforce running tests regularly. [TODO pic that isn't obviously nicked]

Chapter 3

Implementation

The project essentially consists of two parts; the backend and the user interface. The main flow of control within the program can be viewed as a pipeline, with each stage depending on results from the previous stage, and passing results to the next stage. The user interface provides new input, and either runs through the entire pipeline in the simple case, or updates individual stages in the more complex case. In either case, changes to earlier stages necessitates updating all later stages. The project is further divided into the following parts:

1. **Parser**

The parser extracts relevant parameters from a given shader, and their associated types, precisions and arity. Using annotations extra information for the User Interface is also extracted.

2. **Pipeline specification**

The user-specified shader graph is extracted and converted into a usable format, with dummy shaders as presented to the user being converted into actual shader instances.

3. **Parameter UI generation**

An appropriate user interface for specifying parameters for each shader instance in the pipeline is generated based on the pipeline specification and the shader parameters extracted by the parser.

4. **Pipeline generation**

A list of shaders is generated from the shader tree such that the shaders can be rendered in sequence. This creates GLOW shader objects for each shader and assigns framebuffer objects as necessary.

5. **Rendering**

This stage traverses the pipeline and renders each shader in turn. The resulting pipeline output is then displayed in the user interface.

6. **User interface**

The user interface connects the above stages together in an event-driven manner, and allows the user to input shaders and their connections and parameters. It also provides feedback for errors.

3.1 GLSL Parser

The main job of the parser is to take a shader program and provide an array of parameter names and their corresponding data types, which must be supplied to a shader program for it to run. However the entire GLSL program is parsed, to provide extra information, for example for syntax highlighting. For each parameter, the parser returns an object specifying:

- A parameter qualifier: this can be `uniform`, `varying` or `attribute`. We are primarily concerned with `uniform` parameters since these are exposed to the user
- A type: this may be a fundamental type like `int` or a structure, in which case the type is an object like a parameter object, barring the parameter qualifier
- Precision: either `high`, `medium` or `low`. Unused, but could provide error information
- Arity: if the parameter is an array, the size of the array (this is required to be constant)
- Annotations: these include range and colour annotations

All of these other than the type may be left unspecified. The parser could also return information for the syntax highlighter, but this was not implemented. Since WebGL is based on a very specific version of OpenGL, OpenGL ES 2.0, the parser only needs to support GLSL ES 1.0.17 (author?) [14](author?) [8]. This simplifies the construction of the parser, as we need not worry about different GLSL versions.

Since GLSL's grammar is mostly LALR, I deemed it sensible to use a parser generator rather than expend [TODO yes this is a word] effort writing a parser by hand. Initially I chose to use JS/CC (author?) [4], since it seemed well-documented and quite popular. However, once I had used it to generate an appropriate parser for GLSL, it quickly became apparent that the parser was too inefficient to use for such a complicated grammar. In fact, I was unable to obtain any timing for the parser, as under Google Chrome, even with the simplest inputs, it would fail to produce any output before Chrome itself decided to kill the process. While it may have been possible to obtain timing information, this is clearly unacceptable.

I then found Jison (**author?**) [3] — a JavaScript port of GNU Bison. Bison is a venerable parser generator which is part of the GNU project, but it is only capable of generating parsers in C, C++ and Java. Jison was able to parse even quite complicated input in a reasonable amount of time, as is discussed in my evaluation.

As an alternative to using Jison, which is relatively untested compared to Bison, I could have used server-side parsing. However, this would have imposed the need for websites using my project to compile code to run on their servers, and configure their web servers appropriately. Conversely using Jison, no server-side processing is required at all, much reducing the barrier to using my project.

3.1.1 Grammar conversion

The grammar for GLSL is provided for GLSL ES in the specification (**author?**) [14], in BackusNaur Form (BNF). Since Jison's specification language is roughly based on BNF, the conversion is mostly straightforward as shown in Listing 3. The additional annotations between curly braces are constructing parameter type objects to be passed upwards. Token specification is likewise mostly straightforward, with the exception of struct-related tokens as discussed below, and integer / floating-point literals which must be represented via regular expressions. Care must also be taken with the ordering of token definitions, in particular the IDENTIFIER token is capable of capturing most keywords and must be placed after them.

3.1.2 Preprocessor

GLSL contains a preprocessor language much like that used in C. It is very restricted and only allows a few simple substitutions and if/else statements, error messages and compiler information that can be ignored according to the specification. It is often used to make decisions depending on the GLSL language version, but since WebGL currently only supports one version of GLSL it is not generally used in WebGL programs. Implementation would require a simple initial pass over the shader code before forwarding it to the main parser. Because of the limited utility of the preprocessor, I decided not to implement it in my parser, but this would be an easy extension.

Listing 3 Translation of GLSL grammar rules.

//Input: BNF form from specification

```
type_specifier:
    type_specifier_no_prec
    precision_qualifier type_specifier_no_prec
```

//Output: Jison version

```
type_specifier:
    type_specifier_no_prec {
        $$ = {type:$1};
    }
    | precision_qualifier type_specifier_no_prec {
        $$ = {type:$2,prec:$1};
    }
    ;
```

3.1.3 Annotations

In addition to parsing raw GLSL, the parser was intended to be able to detect additional annotations to parameters. These annotations would provide extra information to enrich the usability of the parameter UI, for example by specifying the expected range of values taken by a parameter or, in the case of vectors, whether to provide a colour picking interface. In order to keep compatibility with GLSL, I decided to provide these annotations via comments, as shown in Listing ??.

According to the GLSL specification (**author?**) [14], comments are not to be handled by the parser, but rather be stripped out by a preprocessor. Therefore annotations are achieved by way of my own preprocessor that translates GLSL into ‘annotated GLSL’ by replacing comments of a certain form with code using novel keywords, an example of which we can see in Listing 4. Note that the introduced keywords are named to avoid conflicting with likely actual variable names. Other types of comment are simply removed, as required of the preprocessor by the specification. Since the syntax of these annotations is quite simple, this can be done using regular expressions. This annotated GLSL can then be passed to the parser, the grammar of which is modified to accept these annotations. An example of the rule changes can be seen in 5.

Listing 4 Transformation to annotated GLSL.

//Standard GLSL

```
uniform highp float y; //range 0,100
uniform lowp vec3 c;    //colour
```

//Annotated GLSL

```
uniform highp float y __range__annotation 0,100;
uniform lowp vec3 c __colour__annotation;
```

3.1.4 Struct parsing

The GLSL Specification (**author?**) [14] allows a shader program to define and use new types of structures. These structures behave in essentially the same way as structs in C, subject to certain restrictions concerning self-references. Structure names and field names are subject to the same restrictions as normal identifiers. Since this means that a token may be considered an identifier or a type or a field identifier depending on its usage, this makes parsing GLSL as parsed by for example Google Chrome technically context-sensitive.

However, the parser used, Jison, only supports LALR grammars. General context-free parsers are notoriously complicated and tend to have high time complexity. Given that the grammar is otherwise LALR, I deemed these algorithms to be far too inefficient.

The grammar as provided in the specification introduces extra tokens for structure and field names, skipping over the problem of context sensitivity entirely, but suggesting that the lexer should be able to decide. Therefore by way of an initial, straightforward implementation of struct parsing, I chose to make the parser simply append structure/field names to the appropriate part of the matching logic of the lexer whenever a new structure definition was encountered. This will fail to parse certain shader programs. for example ones where the name of a structure is also used as the name for a variable, or more likely, some field sharing the name of a variable. Since these programs are reasonably unusual, this was deemed to be acceptable for this project.

Listing 5 Modifications to grammar rules.

//Original rule

```
init_declarator_list:
    single_declaration
    | /* omitted */
    ;
```

//New rules

```
init_declarator_list_unannotated:
    single_declaration
    | /* omitted */
    ;

init_declarator_list:
    init_declarator_list_unannotated
    | init_declarator_list_unannotated COLOUR {
        for (i in $$) { $$[i].colour = true; }
    }
    | init_declarator_list_unannotated RANGE INTCONSTANT COMMA
    INTCONSTANT {
        for (i in $$) { $$[i].range = {bottom:$3,top:$5}; }
    }
    ;
```

3.2 Pipeline specification

The purpose of the pipeline specification stage is to take input from the user specifying the shaders to be used and the connections between them — that is, when a shader uses the framebuffer another shader has drawn to as an input texture. Note that at this point the terminology used is similar to that used by the GLOW project rather than the general OpenGL definition, in that the term ‘shader’ is used to refer to the combination of a vertex and a fragment shader. While this may seem unorthodox, it is not generally useful to talk about individual vertex and fragment shaders in this project. Where the distinction is unclear, the term ‘shader program’ is used to refer to an individual GLSL program.

For the actual pipeline specification, I could either develop my own specification language or reuse some existing language. Since this project uses WebGL, I decided to use JavaScript. This has the added bonus of enabling the user to specify a more dynamic pipeline. This is particularly useful given the slight variations in WebGL implementation between web browsers. For example, the pipeline could choose to use a slightly simpler version of a shader for Firefox, given the worse performance on that browser (in my personal experience), or perform some precomputation of values. This would also make the addition of some facility for animation simpler, since this could use some arbitrary JavaScript fragment executed between each render loop, using the same interface to change parameter values.

While I could have required pipeline specifications to produce an actual shader graph in its entirety, the objects used internally by my project are sufficiently complicated that I did not want to expose them directly to the user. Therefore, a simplified environment containing methods for creating dummy shaders is created. These dummy shaders can then be converted into the internal format later. In the interface as presented to the user, most of the internal complexity of a shader instance is hidden. As far as the user is concerned, an instance has parameters, a name and a size, but nothing else. Using the output of another shader instance uses the same interface as any other parameter. The actual behaviour, in which an FBO is written to by one instance and read from by another, is left implicit.

To construct this environment, the pipeline specification text (as JavaScript) is evaluated within a function containing the appropriate function definitions. These definitions include shorthand functions like `sampler2D()` for specifying

ing parameters (in this case, a texture), and shorthand objects with convenient names like `gaussHShader` in Listing 6 for creating shader dummy objects. These objects merely contain the shader instance name, type — that is, the name of the shader it will be an instance of — the size of the framebuffer object to which it will write and any specified parameters.

The resulting directed acyclic graph (or DAG) of dummy objects is then traversed, creating an appropriate graph of actual shader instances. In the initial implementation, this includes placing parameters into an object as needed by GLOW, attaching a reference to the appropriate shader and adding name and size information. To make later updates simpler, this was later modified such that parameters that reference another shader are considered separate — these must have an actual parameter referencing the correct framebuffer object generated. uniform parameters are further separated from other qualifiers (attribute and varying). Later, this stage was also modified to also include early initialisation of GLOW objects. This is discussed in Section 3.6.

Listing 6 shows a simple example of a pipeline specification. This corresponds to the shader graph in Figure 3.1. The variable “output” is special, in that it corresponds to the highest node in the ‘tree’ and therefore the final node in the pipeline — the shader to be rendered to the screen. Notice that some parameters are left blank — these can be provided later via the parameter UI. Individual shader instances can be given names for identification within the UI on initialisation, and the size of the FBO to which the shader is drawn can also be specified. When these are not specified, an identifier is generated based on the shader name and the size is assumed to be the size of the canvas. Unfortunately, since JavaScript lacks a way to dynamically find the variable names within a given context, it is not possible to automatically infer shader instance names without using a JavaScript parser, which is beyond the scope of this project (and would not always be able to infer names since shader instances can be anonymous).

Listing 6 Example pipeline specification.

```
output = new gaussHShader("output");
gaussv = new gaussVShader();
input = new textureShader("input",{width:800,height:600});

output.img = gaussv;
gaussv.img = input;
```




Figure 3.1: The shader graph for the pipeline specified in Listing 6.

3.3 Parameter UI generation

This stage takes a tree of shaders and creates the appropriate HTML user interface. For each parameter of the shader which needs specifying — that is, is of type `uniform` and is not already specified — it produces an appropriate set of input elements. For each parameter a label based upon either the provided name or a hopefully informative generated name is generated, along with a type-specific label and widget.

In the initial implementation, the appropriate shader values would be updated when a new render was requested. However this is inefficient, and makes the user interface less responsive. Later implementations therefore use a different method: each input element has associated event listeners that will update the value within the shader, correctly invalidate the relevant GLOW cache and request the preview be re-rendered. This is achieved cleanly thanks to the presence of closures in JavaScript: the UI generating function has access to the relevant GLOW object for the current parameters, and as such event handles attached within it also have access to the correct object, without the need for less efficient navigation of the shader graph.

While simple text boxes provide complete control over the parameters, for certain types of parameter more useful widgets can be provided. The most obvious example is for `vec3` and `vec4` (vectors of length 3 and 4) parameters, which are often used to specify colours. For such parameters, I constructed a colour picking widget based on the web colour picker by Stefan Petre (author?) [13], as can be seen in [TODO PIC].

Another addition which would make use of the range parameter annotations (see Section 3.1.3) would be sliders for numeric parameters. Texture parameters assume sensible defaults rather than requiring the user to specify the verbose information required by GLOW. This was skipped for time reasons.

Struct handling

Since GLOW expects structure parameters to be specified by ordinary JavaScript objects, which are also used to specify parameters, structs can be easily and simply handled by recursing on the parameter generation function. While for languages like C, this could produce an infinite loop if the structure references itself, GLSL does not permit this. However, the parser will not

reject such a self-referential structure definition, which would cause infinite recursion. This is avoided using a counter set to a sufficiently large value that no valid GLSL program could reach it. This is possible since GLSL programs are required to compile to a sufficiently small amount of instructions, due to the limitations of stream processors.

3.4 Pipeline generation

Pipeline linearisation

At this stage in the process, a directed, acyclic graph (DAG) of shader instances has been generated, with connections representing dependencies between instances — that is, instances that write to textures that will be read by other instances. In order to actually render the DAG, it must be linearised it such that the node labeled “output” is last. This is the primary task of the pipeline generation stage, namely generating a list containing the nodes of the shader DAG. Since JavaScript uses pass-by-reference for objects, we can retain the original DAG at no extra cost. [TODO example] [TODO add list to check for cycles]

This linearisation is achieved using a topological sort. The algorithm described by Cormen et al.[TODO] (**author?**) [9] was used. The algorithm is shown in Listing 1.

Algorithm 1 Pseudocode for topological sort algorithm

```
L ← Empty list that will contain the sorted nodes
S ← Set of all nodes with no outgoing edges
for all node n in S do visit(n)
end for
function VISIT(node)
    if n has not been visited yet then mark n as visited
        for all node m with an edge from m to n do visit(m)
        end for add n to L
    end if
end function
```

Note that there is only one node in *S*, the output node, thus the initial for loop can be replaced with a single call. Additionally, all nodes must be marked as not visited before running this algorithm, as the linearisation procedure may be run multiple times. However, the algorithm shown in Listing 1 cannot detect when the graph contains a loop. This situation will only occur if the user inputs a shader graph containing cycles — in such a case, the program should produce an error.

Pipeline initialisation

In this stage, each shader object in the pipeline is assigned an actual GLOW shader. This is the object that contains the actual compiled shader that will be called at rendertime. Initially, this was done in the most straightforward manner possible, by simply iterating through the pipeline and generating new GLOW objects after any modification. All parameter values would be retrieved from the UI at this point. This required the user to manually request that the pipeline is updated before rendering.

Recreating all GLOW objects and reading in all parameter values after a single parameter change is very inefficient, especially since this process would run after every individual parameter change. Therefore this stage is avoided where possible by ensuring that changing parameters does not require a complete update — GLOW objects are created before the parameter UI generation stage, with dummy parameter values, and pipeline initialisation is only run after a change to the pipeline (as discussed in Section 3.6).

3.5 Rendering

The rendering process itself is kept simple, since I offloaded as much work as possible to other stages. This is important for later extensions involving animation, where we want the render loop to run as quickly as possible. The render loop simply iterates through the shader pipeline list, binding the framebuffer object associated with each shader, rendering the shader and then unbinding the FBO. The FBO binding is skipped for the final shader, which is instead rendered to the preview context. While the initial render function cleared the GLOW cache for each shader instance at the start of each call, we can avoid this as discussed in Section 3.6.4.

3.6 Optimisations

As stated in the introduction, one of the main aims of this project is to make the User Interface as interactive as possible. To be more specific, the delay between making some change and seeing the effect of this change should be as short as possible. Since changes on this timescale will usually be quite small, for example changing a single parameter, or modifying a single shader, by reusing existing state where possible this delay can be significantly reduced.

3.6.1 GLOW early initialization

As discussed in Section 3.3 and Section 3.4, in the initial implementation all stages following Parameter UI Generation must be called, at the users request, following each change to a parameter. This includes a step in which every parameter in the UI is updated, initialisation of GLOW shader objects, and pipeline linearisation. This is obviously suboptimal, both since all parameters must be updated and redundant steps be re-run, and since the user is required to request re-rendering, which reduces the level of interactivity.

The solution to this problem is briefly mentioned in Section 3.4. Rather than creating GLOW objects — and therefore incurring compilation overhead — after all parameters have been specified, we substitute dummy parameters of the appropriate type, and create the GLOW objects during shader tree creation (that is, during the pipeline specification stage). Dummy parameters are obtained by modifying the parameter UI generation stage to generate a dummy value for each parameter as well as the actual UI. Parameter changes now generate events which update the relevant GLOW object's parameters, and request a new render. The entire pipeline initialisation stage can now be skipped. Updates also invalidate the relevant GLOW cache as discussed below.

3.6.2 Modified shader detection

The implementation as discussed above performs the entire shader parsing through pipeline initialization process, throwing away all previous data, whenever a shader program is modified. Since only one shader program can be modified at once, this is sub-optimal. This is particularly important since we would like to be able to render a new preview of the pipeline output as

quickly as possible. The most obvious optimisation here is to only update shader instances in the pipeline that use the modified shader. This is achieved by tagging shader instances with the name of their shader. Shader update calls then only need to attach new GLOW shaders to the appropriate shaders, using their existing parameters. Note that this is only applicable if the required shader parameters have not been added to or changed, since all parameters must be specified. In this case the old behaviour can be falled back on.

3.6.3 Parameter value propagation

Some of the advantages in interactivity introduced by the optimisations made in Section 3.6.2 will be useless, since parameters specified for the modified shader will be lost, and will need to be re-entered by the user. However we cannot simply reuse the previous shader parameters since the parameters taken by said shader may have changed. The set of parameters for the modified shader must first be compared to the new required parameters, and dummy parameters generated where appropriate and any extra parameters must be discarded (these extra parameters are ignored, and we could try to keep these parameters in case they are used in the future [TODO: if I get really bored]).

3.6.4 GLOW cache

For efficiency reasons, GLOW itself tries to cache as much as possible in GPU-accessible memory. In particular, parameter values are cached and updates to parameters must respect this and invalidate the cache before the next render call. In the initial implementation, the GLOW cache was cleared at the start of each render call. This is inefficient, and was soon modified such that the GLOW cache was only cleared after parameter changes and pipeline changes. However, this is still not particularly efficient, as these gains are mostly only noticable during animation. [TODO: can do better?]

3.6.5 FBO reuse

In the initial implementation, FBOs are simply discarded and new FBOs are allocated after each pipeline change. This is time consuming and could be

avoided. However, since FBOs come in different sizes, we cannot naively allocate some 'pool' of available FBOs. This is achieved by simply keeping separate pools of already allocated FBOs for each size of FBO that has been used in the past. Shader updates using the event-based system described above can reuse the existing FBO for each instance, this is an additional benefit of the optimisations discussed.

3.7 User Interface

3.7.1 Editor

As discussed in 2.3, I chose to use the CodeMirror text editor to provide GLSL syntax highlighting to users. While CodeMirror does not provide GLSL highlighting, it does provide a general, 'C-like' highlighting mode, and simple hooks for custom parsers. Initially, I based my syntax highlighting on the 'C-like' mode provided. This was a simple matter of providing the correct keywords. However, this only provides simple highlighting. Given that each shader is parsed anyway, it would seem sensible to modify the parser used for parameter extraction to also provide syntax highlighting information. A particularly common feature provided by IDEs is highlighting the line on which a syntax error occurs. [TODO: actually do this]

3.7.2 Layout

The initial layout of the UI was straightforward, consisting of a pair of text editors for each shader program, a text editor for the pipeline and parameter specification, the parameter UI, a sequence of control buttons and a preview box. This is shown in Figure 3.2. While this layout was spartan and not user friendly, it was sufficient for initial testing, and allowed for easy debugging of separate stages.

3.7.3 Improved Layout

The layout described above suffers from a number of usability problems. Firstly, the UI does not fit at all on an average-sized screen. This requires

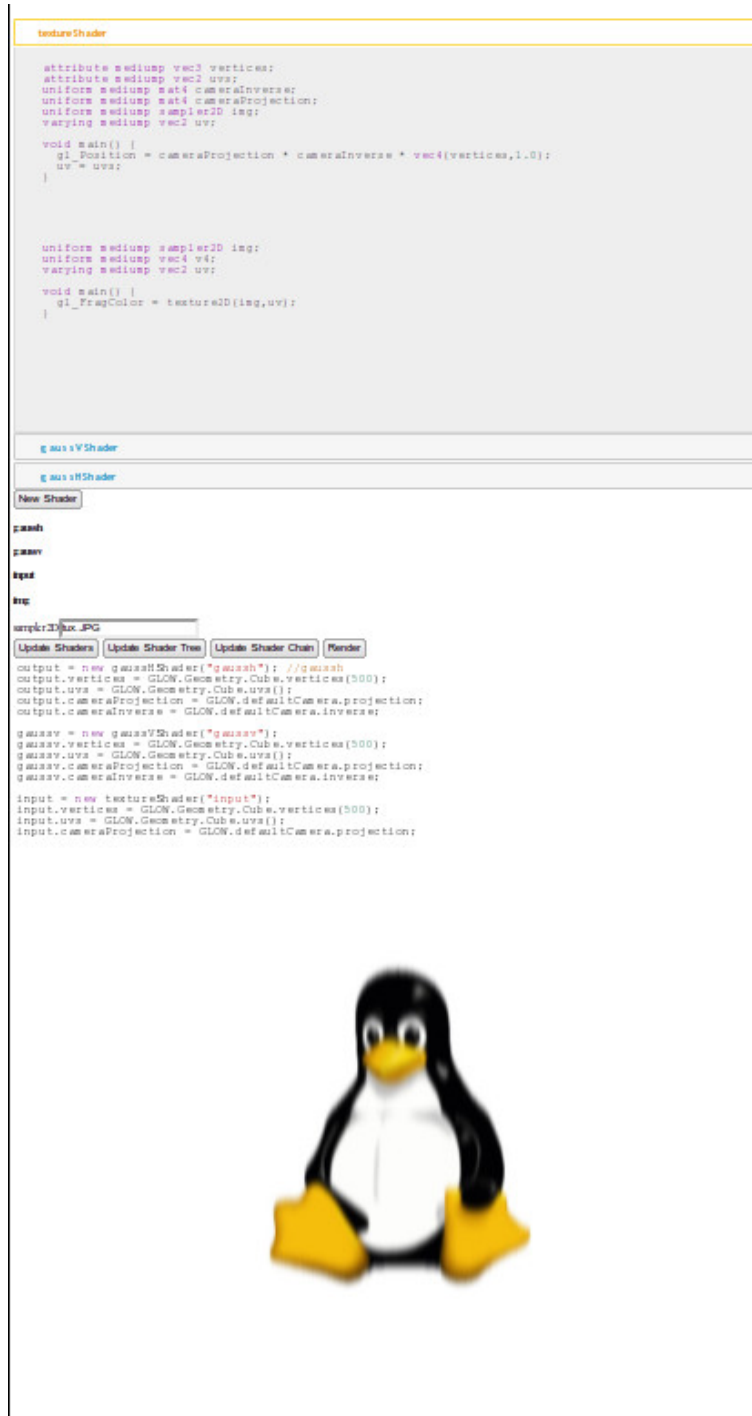


Figure 3.2: Original Layout.

the user to be constantly scrolling to use it, adding a delay between the user making a change, and seeing its effect. Secondly, the control buttons are un-intuitive, and the need for the user to use them whenever they wish to see the effect of a change further increases the expected length of a user's development cycle. A new layout was designed, with inspiration taken from the typical construction of IDEs, where tiling is employed to maximize screen usage. From [TODO another picture] we can see the improved layout. Shader programs occupy the top left side of the screen, and can be switched between by clicking on the shader's name. The right hand side of the screen is occupied by the parameter specification UI at the bottom, and the preview at the top. The control buttons have been eliminated entirely using the hooks developed in Section 3.6. Compared to the old layout, this provides:

- Minimal scrolling — shaders/pipeline specification/parameters can be changed without having to scroll to see the result
- Improved interactivity — updated render provided as soon as possible, rather than on request
- Improved usability — no confusing buttons to press requiring knowledge of internal program workings
- Error feedback — introduction of an error popup means users are informed of errors, but screen space is not wasted otherwise

Chapter 4

Evaluation

Evaluation of this project chiefly consists of two parts, correctness and performance. Correctness evaluation is intended to demonstrate correct functioning of the core project. This consists of two parts: parser testing, to ensure parameters are correctly provided to the user interface, and output testing, to verify that the output of shader pipelines is visually as expected. Performance evaluation provides an analysis of the effect in terms of both performance and resource usage of some of the optimisations. In particular, the effect of avoiding unnecessary shader recompilation and later steps on performance is evaluated.

4.1 Parser correctness

In order to assist with development of the parser, and in particular to assist in regression testing, a test harness was developed for the testing the correctness of parameters returned by the parser. Since the parser does not access any browser-specific APIs, it was possible to perform this testing in a simple, automated way on the command line. Each test is specified by a shader program, and the result that should be returned by the parser — the parameters and their associated information, any struct definitions, and whether the parser should fail to parse it and raise an exception, in the case of an invalid program. Parser output for each test specification is then compared against the actual output for a set of tests. These include correctly throwing/not raising an exception, returning the correct parameter names/a correct subset of the parameter names, all/some of the correct parameter types, the correct struct

definitions, and so on. An additional script was written to perform these tests over all git revisions.

In Listing 7 we see a small sample of the complete test results for an individual version, verbatim from actual command line output. The second column lists individual tests — here, we see some very basic tests verifying that for example `uniform int` is parsed correctly. The third column lists the sub-tests for each test: these are identical for each test, with the exception of tests that are intended to fail to parse. The right column shows whether the sub-test is passed. We can see that for the varying `lowp vec4`, the parameter types are incorrect for this git revision. Listing 8 shows a sample of the test summary for some of the later git revisions. Each line represents a different git revision. Note that there are two sub-tests consistently failed: these correspond to tests testing correct struct parsing for structs sharing names with variables or fields, which cannot be parsed with the current parser implementation, as discussed in Section 3.1.4.

Listing 7 Sample of verbose test results

| Module | Test | Assertion | Result |
|--------|--------|------------------------------|--------|
| glsl | uint | Successfully parses | ok |
| | | One correct parameter | ok |
| | | Correct parameter names | ok |
| | | Correct parameter names only | ok |
| | | Correct parameters | ok |
| | | Correct structs | ok |
| | ufloat | Successfully parses | ok |
| | | One correct parameter | ok |
| | | Correct parameter names | ok |
| | | Correct parameter names only | ok |
| | | Correct parameters | ok |
| | | Correct structs | ok |
| | v1vec4 | Successfully parses | ok |
| | | One correct parameter | ok |
| | | Correct parameter names | ok |
| | | Correct parameter names only | fail |
| | | Correct parameters | fail |
| | | Correct structs | ok |




Figure 4.1: Percentage of parser tests passed vs. git revision.

Listing 8 Sample of test summary

Summary:

| File | Failed | Passed | Total | Runtime |
|--|--------|--------|-------|---------|
| /home/joseph/ws/clone/shomp/parser/gls1.js | 2 | 34 | 36 | 177 |
| /home/joseph/ws/clone/shomp/parser/gls1.js | 2 | 34 | 36 | 179 |
| /home/joseph/ws/clone/shomp/parser/gls1.js | 2 | 34 | 36 | 179 |
| /home/joseph/ws/clone/shomp/parser/gls1.js | 2 | 34 | 36 | 173 |
| /home/joseph/ws/clone/shomp/parser/gls1.js | 2 | 34 | 36 | 175 |
| /home/joseph/ws/clone/shomp/parser/gls1.js | 2 | 34 | 36 | 178 |

In Section 4.1 we see a graph showing the percentage of tests correctly completed, from the first revision including a parser. The notable jumps in tests passed correspond to implementation of the currently used parameter format, scoping corrections (i.e. discarding parameters hidden within functions) and structure support. Not all tests are passed at the final version — these tests are those that contain structure definitions and variables with the same name, the parsing of which is not possible with my current parser, as discussed in Section 3.1.4. Also tests that use preprocessor commands will fail since a pre-processor was not implemented.

4.2 Pipeline correctness

In order to demonstrate visually correct shader pipeline output, a small number of example pipelines were constructed. This was achieved by taking a sample of simple WebGL example programs found online, and modifying them to work within my project. Output from the modified versions could then be compared with the originals. Ideal example programs would use a few different shaders connected together to produce a single visual effect. Below we can see samples of the output of the original programs alongside my reimplementations.

Gaussian blur

Gaussian blurring of an image can be naturally divided into two stages, a horizontal and a vertical blur. A simple pipeline is shown in Listing 9, and its graph in Figure 4.2. Parameter values have been mostly omitted in the pipeline for clarity. Each stage performs the function $G(x) = \frac{1}{\sqrt{2\pi\sigma^2}}e^{-\frac{x^2}{2\sigma^2}}$ in one dimension, which is equivalent to the function $G(x, y) = \frac{1}{2\pi\sigma^2}e^{-\frac{x^2+y^2}{2\sigma^2}}$. The first shader, “input” is strictly unnecessary. Figure 4.3 shows a comparison of the original to the blurred image.

Listing 9 Gaussian blur pipeline.

```
output = new gaussHShader("output");
gaussv = new gaussVShader("gaussv");
input = new textureShader("input",{width:800,height:600});

output.img = gaussv;
gaussv.img = input;
```

Sanity test

This example was developed in order to ensure that long pipelines are rendered correctly. It uses a long chain of only one type of shader. These shaders merely increment the value of the red channel, so a chain of length n should increment the initial red value by n times the step size. A simplified version




Figure 4.2: Corresponding shader graph for Listing9.




Figure 4.3: Comparison of original to blurred image.

of the pipeline is shown in Listing 10. This example demonstrates how using an existing programming language for pipeline specification allows concise construction of more complex pipelines.

Listing 10 Sanity test pipeline.

```
output = new step();
var c = output;
for (var i=0;i<10;i++) {
    c.img = new step();
    c = c.img;
}
c.img = sampler2D('blank.JPG');
```

4.3 Performance

Effect of modified shader detection

In Section 3.6, we discussed the modification of the pipeline initialization process to avoid unnecessary shader re-parsing, compilation, and pipeline generation after modification of a shader. In order to compare the performance of the updating procedure before and after these optimisations were introduced, a custom test harness was developed. Since the existing test harnesses for JavaScript that provide performance testing are mostly quite heavyweight, being intended for large-scale website testing, I implemented this from scratch.

Each test specification consists of a set of shaders, and a shader tree specification in which all necessary parameters are specified. Once a shader pipeline has been generated for a given specification, a shader update event is generated for a particular shader, and its completion timed. This is repeated 10 times to obtain an average and range. Similarly the completion of a complete re-generation of the entire shader pipeline is timed 10 times. Since WebGL is intelligent enough to avoid recompiling identical shader text, shaders programs are padded with a different amount of trailing whitespace between iterations. This is sufficient to defeat the caching behaviour without changing the semantic meaning of the shader program. This is repeated for each shader in the test specification. From Figure 4.4 we can see the result of these tests for a particular test specification. It is clear from this graph that these optimisations have a significant effect on the speed of an update. While for smaller

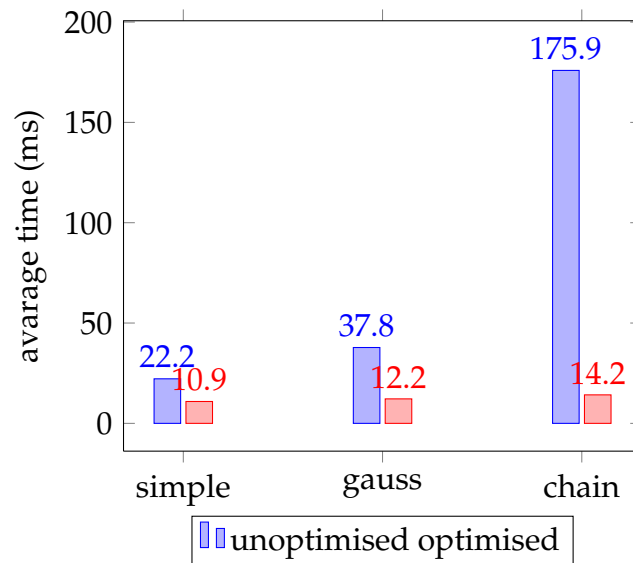


Figure 4.4: Graph comparing shader update times pre- and post-optimisation.

examples like “simple”, which consists only a single shader, the difference in performance is unlikely to be noticed, for a more complicated example the benefit is clear.

Chapter 5

Conclusion

5.1 Main Results

The project has been successful. The core of the project has been completed and functions correctly, and with the addition of the various optimisations and extensions, is sufficiently interactive to be useful. I have made the project source code available to others via GitHub, and I hope that it will prove useful. In Table 5.1 we can see that the core criteria were completed, and from Table 5.2 we can see that a significant number of the extensions were completed. Of the extensions not implemented, two of them were deemed unsuitable. Texture unification as originally envisaged is supported trivially by the pipeline specification stage, and code separation would be sufficiently complex to be outside the scope of this project, and is unlikely to be worth the extra effort. In addition, a number of optimisations were made not listed in the table below.

5.2 Lessons learned

The combination of WebGL and Linux is still a little unstable at times, and as such I would be less inclined to rely so much on a very recent technology for future projects.

While I was initially apprehensive about using JavaScript given its reputation as an ugly language to work with and my experience of Internet Explorer 6-era JavaScript, I was pleasantly surprised by its novel object model and quite functional underpinnings. Particularly useful was the ability to apply a class

| Criterion | Successful? |
|--|---------------|
| Interface to construct shaders, specify connections between shaders | Yes |
| Detection of parameters of set of shaders, parameters presented to the user somehow | Yes |
| Sample scene output for single shader | Yes |
| Demonstration of correct composition (with sample scene output) for various shaders with pipeline specifications | No [SEE EVAL] |
| Provide basic syntax highlighting of GLSL in interface | Yes |

Table 5.1: Core Criteria

| Criterion | Successful? |
|---|------------------|
| Reuse of FBOs in shader pipeline | Yes |
| Automatic unification of identical textures | No ¹ |
| Splitting shared code off into separate shaders | No ² |
| Avoiding recompiling shaders unnecessarily | Yes |
| Presentation of sliders with range detected from range annotations in shaders | Yes[TODO broken] |
| Some facility for animation | No |
| Shader and/or scene test suite | No |

Table 5.2: Extensions

constructor to an already existing object, and the ability to treat objects much like associative arrays. Combined with a good set of libraries like jQuery, I found modern JavaScript about as pleasant to work with as languages like Python or Ruby.

5.3 Further work

As this project was only partially focused on the development of a user interface, the current UI is fairly basic. Further work could be done to make the UI more usable, and possibly to conduct a user study to test its effectiveness. More sophisticated widgets and better error feedback would be a good start. The extensions listed in Table 5.2 were dropped for time reasons, but could also be incorporated with not too much additional effort. A number of techniques are used for dealing with context sensitivity in C without using a general context-sensitive parsing algorithm, which could potentially be applied to parsing GLSL with structs correctly.

Both the specification and implementation of WebGL have changed significantly since their first proposal, and I expect them to continue to do so for some time yet. As such further work may be necessary in the future as these change. Of particular note is the current lack of multiple render targets in current implementations (**author?**) [7]. If in the future implementations begin to support multiple render targets, adding support for these to my project will be a moderately non-trivial task.

Bibliography

- [1] Codemirror. <http://codemirror.net/>.
- [2] I am GLOW. <http://i-am-glow.com/>.
- [3] Jison. <http://zaach.github.com/jison/>.
- [4] JS/CC parser generator project homepage. <http://jscs.jmksf.com/>.
- [5] three.js. <https://github.com/mrdoob/three.js/>.
- [6] V8 JavaScript engine. <http://code.google.com/p/v8/>.
- [7] WebGL and HTML5 challenges for the future. <http://codeflow.org/entries/2011/sep/11/webgl-and-html5-challenges-for-the-future/>.
- [8] WebGL specification. <http://www.khronos.org/registry/webgl/specs/latest>.
- [9] Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. *Introduction to Algorithms*, chapter 22.4: Topological sort, pages 549–552. The MIT Press, 2 edition, 2001.
- [10] M. Haverbeke. *Eloquent JavaScript: A Modern Introduction to Programming*. No Starch Press Series. No Starch Press, 2011.
- [11] David Miller. Setting up emacs as a javascript editing environment for fun and profit. <http://blog.deadpansincerity.com/2011/05/setting-up-emacs-as-a-javascript-editing-environment-for-fun-and-profit/>, May 2011.
- [12] OpenGL, D. Shreiner, M. Woo, J. Neider, and T. Davis. *OpenGL Programming Guide : The Official Guide to Learning OpenGL(R), Version 2 (5th Edition)*. Addison-Wesley Professional, August 2005.

- [13] Stefan Petre. ColorPicker – jQuery plugin. <http://www.eyekon.ro/colorpicker/>.
- [14] Robert J. Simpson. *The OpenGL ES Shading Language*. 2009.
- [15] Sjoerd Visscher. Higher order programming. <http://w3future.com/html/stories/hop.xml>, October 2005.
- [16] Felix Woitzel. WebGL shader lab reactor-diffusion mix. <http://cake23.de/diffusion-mix.html>.

Appendix A

Project Proposal

Part II Computer Science Project Proposal

Shader Compositor

Joseph Seaton, Fitzwilliam College

Originator: Christian Richardt

21 November 2011

Special Resources Required

None

Project Supervisor: Christian Richardt

Director of Studies: Dr R. Harle

Project Overseers: Dr S. Holden, Dr D. J. Greaves

Introduction

Designing shaders can be laborious work, involving endless back and forth between tweaking the code (often minor visual tweaks) and examining the visual output. While programs exist to provide previews of individual shaders, modern software, especially for computer game engines, often involves multiple shaders chained together, and little software exists to aid shader writers in testing such shader pipelines. Furthermore for beginners, the initial process of learning of how to use shaders is complicated by this extra legwork to compose shaders. The aim of this project is to alleviate these problems to some degree.

Aside on Shaders

A 'shader' is a piece of code that runs directly on a GPU. Shaders work on the data flow model, in that once the code is set up, many data may be fed through it. In OpenGL, which this project will focus on, such shaders are divided into a number of types. The two types considered here are vertex shaders and fragment shaders although there are others including geometry shaders which could be considered as an extension. Vertex shaders operate on a vertex, and its associated data, performing such operations as transforming a vertex to the screen coordinates. Fragment shaders (approximately speaking) operate on individual pixels, and are often used for operations such as texturing a model. Vertex shaders may pass information to fragment shaders, and the underlying program using the shaders may pass information to both.

The Project

This project intends to provide an interface to allow a user to create a set of shaders, written in a slightly annotated version of GLSL. The attributes of these shaders will be detected, along with the type of each attribute, and, given a specification of how to connect these attributes (provided either via a GUI or a simple language), a pipeline of these shaders will be constructed. It should be noted that the construction of said pipeline will require a number of stages, namely construction of a simple GLSL parser, generation of a DAG of shaders, linearisation of said DAG, and finally application of this DAG to FrameBuffer Objects for offscreen rendering, finishing in some rendered model. Since FBOs

can be reused, an efficient mapping is non-trivial. Furthermore, there is room for optimisation of this DAG consider unification of identical textures, or extraction of common code. The interface should also provide some facility to allow the user to easily modify shader parameters by e.g. selecting textures, or setting the colour of a vector. Addition of simple annotations to the GLSL may be useful here, such as specifying the range of a vector. Since this project is intended to be used partly as a learning tool, I propose to use WebGL, a standard for using OpenGL (including GLSL) on the web, via JavaScript. Since WebGL is cross platform by design and requires only a recent web browser to use, As WebGL is quite a new technology, it is possible that this may turn out to be infeasible, in which case I would use Python.

Resources Required

Since this project will involve OpenGL shaders, a recent graphics card supporting a recent version of OpenGL will be required. To this end I would use my own computer with its AMD HD5450 graphics card. Depending upon the speed at which my test harness runs, I may consider upgrading to a more powerful card such as the HD6770.

Starting Point

I have some existing knowledge of OpenGL and GLSL from personal experience, but only to the extent of small personal projects. I have a reasonable amount of experience with JavaScript having used it in the past for web development work but in conjunction with a server-side language. Substance and Structure of the Project The project involves writing software providing a user interface in which to input a set of shaders and the connections between them and view a sample output of the generated shader pipeline on a sample model. Most notably to do this the project requires software to be written that given a set of shaders and a specification of how to connect them, can parse each shader, extract its parameters, and construct the specified pipeline. Furthermore the user interface should provide the ability to easily vary the parameters of each shader on the fly. This and the above comprise the core of the project. By way of extensions, there is much scope for fleshing out the interface into a fully-featured development tool, with such additions as example shaders, a more complete set of sample scenes, and so on. There is also some room for

optimisation of the composition process, for example the unification of identical textures or extraction of common code between shaders. For evaluation, in order to demonstrate correct composition, a testing framework will be written that takes a set of semantically obvious or simple shaders and composites them in arbitrary ways and tests whether the output matches the expected result. Also pipeline speed/memory use comparisons of before/after the implementation of the extensions will be performed, along with comparisons of recompilation speed before/after optimisation (i.e. detection of modified shader). Further comparisons include comparisons of a generated pipeline against a hand-coded shader, or a comparison of the time complexity of a given pipeline to a JavaScript program carrying out the same task.

Success Criteria

The following should be achieved:

1. Interface to construct shaders, specify connections between shaders
2. Detection of parameters of set of shaders, parameters presented to the user somehow
3. Sample scene output for single shader
4. Demonstration of correct composition (with sample scene output) for various shaders with pipeline specifications
5. Provide basic syntax highlighting of GLSL in interface

Extensions: Demonstration of some optimisations:

1. Reuse of FBOs in shader pipeline
2. Automatic unification of identical textures
3. Splitting shared code off into separate shaders
4. Avoiding recompiling shaders unnecessarily

Demonstration of some additional GUI work:

1. Presentation of sliders with range detected from range annotations in shaders,

2. Some facility for animation (e.g. parameters that are automatically incremented)
3. Shader and/or scene test suite

Timetable and Milestones

Weeks 0–2: 22nd October - 11th November

Initial reading. In particular, the OpenGL 'red book'. Setup of barebones GUI, associated libraries. Preparation of automatic backups. Milestone: dummy interface for single shader.

Weeks 3–5: 12th November - 2nd December

Construction and testing of GLSL parser. Further familiarisation with GLSL. Milestone: ability to parse simple GLSL

Weeks 6–8: 3rd December - 23rd December

Write shader DAG generating code, DAG linearisation code. Initial work on application to FBOs. Milestone: generation of linearised shader DAGs

Weeks 9–11: 24th December - 13th January

Complete actual pipeline given some DAG, the project should be able to render the pipeline to a quad. Complete interface such that the pipeline order can be specified, either via a GUI or via some written specification. Basic parser may be sensible. Milestone: ability to render a given shader pipeline

Weeks 12–14: 14th January - 3rd February

Preparation of progress report. Presentation of shader parameters in GUI. Initial work on texture unification and other optimisations. Initial work on test

suite. Milestone: Progress report prepared. Ability to vary shader parameters from the GUI. Ability to unify some identical textures (or reduced FBO or code usage, depending on extension). Some tests written.

Weeks 15–17: 4th February - 24th February

Presentation of progress report. Further testing of core functionality, expansion of test cases. Bug fixing. Possibly initial optimisation work. Milestone: Progress report given. Large test suite. Lack of obvious bugs in core.

Weeks 18–20: 25th February - 16th March

Additional GUI work and/or optimisation. Slack time for unexpected problems. Optimisations can be skipped if more time is needed for core work. Milestone: Better featured GUI and/or some further optimisations.

Weeks 21–23: 17th March - 6th April

Begin work on dissertation. Programming work should be finished by now, although some expansion of the test suite may be acceptable. Milestone: complete initial sections of dissertation.

Weeks 24–26: 7th April - 27th April

Finish first draft of dissertation.

Weeks 27–29: 28th April - 18th May

Proofreading and submission.